

Proximo Agentic AI Architecture

System Overview

User Input → Orchestrator → Specialized Agents → Response Synthesis → Output

Core Agents

1. Conversation Agent

- **Model:** Fine-tuned Mistral-7B per friend personality
- **Function:** Generate contextual responses
- **Memory:** Vector database of conversation history
- **Tools:** LangChain, MemGPT for long-term memory

2. Emotion Detection Agent

- **Model:** RoBERTa-emotion or BERT-emotion
- **Function:** Analyze user emotional state
- **Input:** Text, voice tone analysis
- **Output:** Emotion classification + intensity

3. Personality Matching Agent

- **Model:** Custom-trained BERT for Big Five traits
- **Function:** Match response style to friend personality
- **Data:** Friend's communication patterns, linguistic markers
- **Output:** Personality-adjusted response parameters

4. Avatar Generation Agent

- **Options:**
 - **Paid:** HeyGen API (\$99/month)
 - **Open Source:** SadTalker + Wav2Lip + Real-Time-Voice-Cloning
- **Function:** Generate talking avatar videos
- **Input:** Text response + friend's appearance data
- **Output:** MP4 video with lip-sync

5. Voice Synthesis Agent

- **Options:**
 - **Paid:** ElevenLabs (\$5-330/month)
 - **Open Source:** Coqui TTS + OpenVoice
- **Function:** Clone friend's voice patterns
- **Input:** Voice samples from friend
- **Output:** Synthesized audio in friend's voice

6. Memory Management Agent

- **Tools:** ChromaDB + custom retrieval system
- **Function:** Store and retrieve conversation context
- **Features:**
 - Long-term personality consistency
 - Context-aware responses
 - Relationship progression tracking

Agent Orchestration Flow

Input Processing

```
python
```

```

class ProximoOrchestrator:
    def process_user_input(self, message, user_id, friend_id):
        # 1. Emotion Detection
        emotion = self.emotion_agent.analyze(message)

        # 2. Retrieve Context
        context = self.memory_agent.get_context(user_id, friend_id)

        # 3. Personality Parameters
        personality = self.personality_agent.get_traits(friend_id)

        # 4. Generate Response
        response = self.conversation_agent.generate(
            message, emotion, context, personality
        )

        # 5. Create Avatar
        avatar_video = self.avatar_agent.create_video(
            response, friend_id
        )

        # 6. Update Memory
        self.memory_agent.store_interaction(
            user_id, friend_id, message, response, emotion
        )

        return avatar_video, response

```

Response Generation Pipeline

User Message → Emotion Analysis → Context Retrieval →
 Personality Loading → Response Generation → Avatar Creation →
 Voice Synthesis → Video Compilation → User Delivery

Technical Implementation

Infrastructure Stack

- **Orchestration:** FastAPI + Celery for async processing
- **Agents:** Individual Docker containers per agent
- **Communication:** Redis for inter-agent messaging

- **Storage:** PostgreSQL + ChromaDB for vectors
- **Caching:** Redis for frequently accessed data

Agent Deployment

```
yaml

services:
  .. orchestrator:
  .... image: proximo/orchestrator
  .... ports: ["8000:8000"]

  .. conversation-agent:
  .... image: proximo/conversation
  .... environment:
  ..... - MODEL_PATH=/models/mistral-friend-{friend_id}
  .....
  .. avatar-agent:
  .... image: proximo/avatar
  .... environment:
  ..... - USE_HEYGEN=${HEYGEN_API_KEY:+true}
  ..... - SADTALKER_PATH=/models/sadtalker
  .....
  .. voice-agent:
  .... image: proximo/voice
  .... environment:
  ..... - ELEVENLABS_KEY=${ELEVENLABS_KEY}
  ..... - COQUI_MODEL_PATH=/models/coqui
```

Cost Optimization Strategy

Hybrid Approach: Open Source + Paid Tiers

- **Free Tier:** Open source tools only
 - SadTalker for avatars
 - Coqui TTS for voice
 - Self-hosted Mistral-7B
 - Cost: ~\$200/month for VPS hosting
- **Premium Tier:** Best-in-class paid tools
 - HeyGen for high-quality avatars
 - ElevenLabs for premium voice cloning

- Cost: ~\$500/month + compute costs

Compute Requirements (Self-Hosted)

- **GPU:** RTX 4090 or A6000 (24GB VRAM)
- **RAM:** 64GB system memory
- **Storage:** 2TB NVMe SSD
- **Estimated Cost:** \$5,000-8,000 hardware + \$300/month electricity

Development Phases

Phase 1: Core Agents (Months 1-3)

- Build conversation agent with Mistral fine-tuning
- Implement basic emotion detection
- Create memory management system
- Simple text-based responses

Phase 2: Avatar Integration (Months 4-5)

- Integrate SadTalker for video generation
- Add voice synthesis with Coqui TTS
- Build friend data processing pipeline
- Basic avatar customization

Phase 3: Premium Features (Months 6-7)

- HeyGen API integration for paid users
- ElevenLabs voice cloning
- Advanced personality modeling
- Real-time conversation capabilities

Phase 4: Optimization & Scale (Months 8-9)

- Multi-agent load balancing
- Response caching and optimization
- Mobile app integration
- Performance monitoring

Key Advantages of Agentic Approach

1. **Modularity:** Easy to upgrade individual components
2. **Scalability:** Each agent can scale independently
3. **Cost Flexibility:** Mix open source and paid tools
4. **Rapid Development:** Leverage existing specialized tools
5. **Quality:** Best-in-class tools for each function
6. **Maintainability:** Cleaner separation of concerns